

TokenBudget: Mapping Where Visual-Token Compression Breaks Edge MLLM Perception

A Technical Report

Model under test: MiniCPM-V 4.6 — a 1.3B edge-oriented multimodal LLM with native `downsample_mode` vision-token compression (4x / 16x).

Budget axis: downsample mode × input resolution → 6 configurations, **82–2,598** effective vision tokens (a 31.7× span).

Dataset: 960 procedurally-rendered perception probes across 8 task families, each with an exact ground-truth answer.

1. Motivation

Edge multimodal LLMs must fit inside a compute envelope. The dominant lever is the *visual-token budget*: how many tokens are spent describing a frame. MiniCPM-V 4.6 exposes `downsample_mode` (4x / 16x) and, through input resolution, a second knob. The model card reports aggregate accuracy under compression, but an application developer faces a harder question:

Which specific perception capabilities survive aggressive token compression, and which collapse first?

"Fewer tokens" and "less perceptive" are not the same thing. Some capabilities degrade gracefully; others fail catastrophically once the budget passes a threshold. This project maps that landscape with a controlled, procedurally-generated perception battery — tasks with known ground truth and a calibrated difficulty axis — and reports where perception *quietly breaks*, and where it does not.

2. Method

2.1 Probe corpus

Images are rendered deterministically with PIL primitives only. Each family varies a single *spatial-frequency / information-density* axis as difficulty rises. Every probe carries an exact ground-truth answer, and the question offers a closed set of answer options so scoring is unambiguous.

family	task	difficulty axis
ocr	read a 5-digit number	font size 96→16 px, additive noise

family	task	difficulty axis
count	count same-color dots	3→24 dots, shrinking radius
mlread	read the 2nd line of a digit column	2→8 lines
colcnt	count red cells in a dense color grid	8→30 cells
spatial	which color is on the left	square size, distractors
color	closest anchor hue (red/orange/pink)	hue offset 40→7
detail	which row of rings has a gap	gap 12→1 px
ringgap	which row of thin rings has a 1px gap	ring radius 60→18 px

20 seeds × 6 difficulties × 8 families = **960 probes**, rendered at 512×512 and downsampled to the configuration resolution at evaluation time.

2.2 Configurations

Six budget points. Token counts are *measured* by the engine itself (image-token count from the processed input), not taken from the model card:

config	downsample	resolution	vision tokens
16x@448	16x	448 px	82
4x@448	4x	448 px	274
16x@672	16x	672 px	347
16x@1280	16x	1280 px	678
4x@672	4x	672 px	1307
4x@1280	4x	1280 px	2598

The budget axis is thus non-uniform: 82 → 274 → 347 → 678 → 1307 → 2598, spanning more than a factor of 30.

2.3 Scoring and analysis

Free-form model output is mapped to the closed answer set by a lenient-but-explicit per-family normalizer (`tokenbudget.check`); a probe is correct iff normalized output equals ground truth. Accuracy per (family, difficulty, config) is reported with Wilson 95% confidence intervals.

A **cliff** is the first budget step (walking from richest to tightest) where accuracy drops ≥ 15 points in one step **and** ≥ 20 points below the richest configuration.

3. Results

All numbers below are computed directly from `data/sweep/sweep.json` by `scripts/paper_facts.py`; the script prints exactly the aggregate curves quoted here. "Aggregate" is the count-weighted mean over the six difficulties at each budget point (n = 120 per point per family).

3.1 The headline: compression is selectively destructive

Per-family aggregate accuracy at the richest and tightest budget points:

family	agg @ richest	agg @ tightest	largest single-step drop	verdict
ringgap	0.99	0.80	0.16 (at 16x@672)	cliff
count	0.46	0.51	0.05	flat, capacity-limited
colcnt	0.29	0.40	0.20	hard, budget-invariant
mlread	1.00	0.98	0.02	robust
ocr	0.99	0.98	0.02	robust
spatial	1.00	0.99	0.01	robust
color	0.97	0.98	0.02	robust
detail	1.00	1.00	0.00	fully robust




Figure 1

Figure 1 — Per-family accuracy vs vision-token budget (log axis, right = richest budget). One line per difficulty level; the red dashed line marks a detected cliff.

The honest reading is more interesting than a clean split:

- **Exactly one task family shows a real, aggregate-level cliff: ringgap**. Everything else either stays within noise of its rich-budget level (ocr, mread, spatial, detail, color) or is hard but invariant to the budget (colcnt, count).
- **count is not budget-fragile in aggregate**. Its wide-looking range in per-difficulty numbers (0.85 down to 0.10) is an artifact of comparing the easiest difficulty at full budget against the hardest at the tightest budget. As an aggregate capability, counting accuracy is nearly flat across the whole budget axis (0.46–0.51). The bottleneck for counting on this model is the model's *capacity*, not the token budget.

- **detail** is fully robust even at 82 tokens, while ringgap collapses — both are "detect a gap" tasks at similar difficulty. The distinguishing property is what the answer requires: detail asks *whether* a coarse gap exists (recoverable from low-frequency content), while ringgap asks *which* of several thin rings carries a 1px gap (requires per-element binding near the model's resolution limit).

3.2 The cliff: ring-gap detection

ringgap is the cleanest demonstration of a compression cliff, and it is non-monotone in a way worth stating plainly:

difficulty	4x@1280 (2598 tok)	16x@672 (347 tok)	16x@448 (82 tok)
d0	1.00	1.00	1.00
d1	1.00	1.00	1.00
d2	1.00	1.00	1.00
d3	1.00	0.65	0.90
d4	1.00	0.50	0.45
d5	0.95	0.10	0.45

Difficulties d0–d2 are *fully robust*: 1.00 at every budget. Difficulties d3–d5 degrade sharply at the 347-token operating point — d5 falls from 0.95 at the richest configuration to 0.10. The aggregate family curve drops 0.867 → 0.708 (15.8 points) across this step, the only aggregate cliff in the battery.

The tightest budget (16x@448 , 82 tokens) shows a *partial recovery* (aggregate 0.80 vs 0.71 at the 347-token point). This is not a "continue falling" curve. The most plausible reading is that at 82 tokens the model can no longer bind individual rings and falls back to a coarse heuristic that is wrong-but-better-than-random on the hardest items. We flag this explicitly because a naive cliff detector reports the 347-token dip as the breakpoint; the full curve shows a V-shaped failure mode rather than a monotone collapse.

3.3 Per-element counting is capacity-limited, not budget-limited

count does *not* erode with budget. The aggregate curve is flat at 0.46–0.51 across the whole axis, and the per-difficulty extremes (0.85 at easiest/full, 0.10 at hardest/tight) reflect difficulty and model capacity, not compression. This is the correct null result: on a 1.3B edge MLLM, counting small dot sets is already near chance at full budget, so compression cannot be blamed for a failure that is present at 2,598 tokens. Any deployment claim about counting under compression must first establish the full-budget baseline.

3.4 Dense-grid color counting is hard, but not budget-limited

colcnt — counting red cells inside a dense color grid — is the hardest task in the battery, but its difficulty does **not** come from compression. The aggregate sits at 0.24–0.48 across the whole axis with no monotone trend (0.29 at the richest configuration, 0.40 at the tightest, peak 0.48 mid-axis), and per-difficulty accuracy

ranges from 0.15 (hardest) to 0.75 (easiest) at every budget. The single largest step drop (0.20) happens *between the two richest mid-axis configurations* (4x@672 → 16x@1280), i.e. across a compression-mode boundary, not toward the tight end of the axis. This is a difficulty/heterogeneity effect, not a budget effect. Like count , this task exceeds what a 1.3B edge MLLM can do reliably even at full budget; it documents a genuine capability boundary of the model class rather than a compression boundary, and no budget-sensitivity claim can be built on it.

3.5 What survives compression

ocr , m1read , spatial , detail , and color are effectively budget-invariant:

- ocr and m1read stay at 0.95–1.00 across all six configurations — reading a digit string or a specific line needs only a coarse feature of each glyph.
- detail is 1.00 everywhere: the gap in a coarse ring is a low spatial- frequency feature that survives 16x downsampling.
- spatial and color stay above 0.90: layout and hue discrimination are dominated by low-frequency content.

The robustness of these tasks is not trivial — it shows that the 16x compression in MiniCPM-V 4.6 does not destroy *all* information; it selectively destroys the high-frequency, per-element information that only one task in this battery (ringgap) actually requires.

4. Discussion

4.1 The compression landscape, in one sentence

For a 1.3B edge MLLM, **token compression destroys per-element binding while global structure survives — but only one of our eight probe tasks actually demands per-element binding**. Fine-gap localization (ringgap) degrades sharply and non-monotonically below ~350 tokens; digit reading, coarse detection, layout, and color survive at 82 tokens. Counting and dense-grid color counting are capacity limits, not budget limits.

4.2 Implications for deployment

- **Choosing 16x vs 4x.** For applications that only read text, recognize coarse objects, or judge layout, 16x@448 (82 tokens) is nearly free: 0.95+ accuracy at a fraction of the compute. For applications that must localize fine structure (e.g. which of several thin elements carries a defect), the budget below ~350 tokens is a reliability cliff; 4x@672 (1307 tokens) or higher is warranted.
- **The cliff, not the slope, is the engineering-relevant number.** A developer wants to know where a capability *breaks*, not just that it degrades. The ringgap cliff at 16x@672 (347 tokens) is the kind of actionable threshold a deployment gate can encode — with the caveat that the failure is V-shaped: at the tightest budget the model reverts to a coarse heuristic that is not monotonically worse.

- **Establish the full-budget baseline before blaming compression.** `count` and `colcnt` look like compression failures until you check their accuracy at 2,598 tokens. Both are already far below the model's capacity ceiling at full budget; compression is not the causal agent. A budget-sensitivity study that skips the richest configuration over-attributes failures to compression.

4.3 Why this looks different from the model card

Model-card accuracy aggregates over natural-image tasks, where compression losses average out and the majority of pixels are background. A probe battery with per-task, per-difficulty resolution exposes the variance the aggregate hides — but it also shows how easy it is to over-read that variance: without the aggregate view, `count` looks like a compression failure when it is in fact a capacity limit.

4.4 Related work

Visual-token reduction is an active design axis for multimodal LLMs. The closest line of work compresses tokens with learned projectors or prunes them during inference: TokenPacker reduces 75–89% of visual tokens at the projector with a coarse-to-fine scheme [4]; FastV prunes low-attention visual tokens in deep layers of an LVLM at inference time [5]; EViT reorganizes inattentive patches in a ViT backbone [6]. A second line studies how input resolution and tiling change perception on high-resolution benchmarks — e.g. LLaVA-UHD's dynamic slicing [7], Qwen2-VL's "native resolution" [8], and MiniCPM-V's own 4x/16x downsample modes [1,2]. What these studies measure is *aggregate* benchmark accuracy after compression. This report complements them with a controlled probe battery that keeps the image content fixed and sweeps only the budget knobs, so the per-capability fate of individual skills (reading, counting, localizing, hue matching) is visible rather than averaged away. It also adds the aggregate-vs-per-difficulty distinction, which we found is the difference between a spurious "clean split" and a truthful "one real cliff" conclusion.

4.5 Limitations

- **Single model, single image scale.** All results are for MiniCPM-V 4.6 at 512×512 rendered probes. The V-shaped `ringgap` curve in particular may be model-specific; we report it as an observation, not a law.
- **n = 20 per (family, difficulty, config) cell, single decode.** Wilson intervals are wide at n=20 (e.g. ± 0.21 at 0.50), and the single greedy decode discards sampling variance. Cells near a cliff should be re-measured with temperature sampling before deployment decisions.
- **Aggregate curves over six difficulty levels can hide per-difficulty behavior.** We report both views explicitly, but a family with heterogeneous difficulties could still mislead if only the aggregate is read.
- **Procedural probes are not natural images.** They isolate specific perceptual skills, but transfer to real scenes is untested.
- **`colcnt` and `count` are hard tasks.** The "budget-invariant" interpretation assumes the tasks are well-formed and that their full-budget level reflects model capacity. Both remain below 0.50 in aggregate even at the richest configuration; their flatness bounds what this battery can say about the budget sensitivity of hard perceptual counting, by design.

5. Reproducibility

```
pip install -e .[gpu,paper,ui]
python scripts/build_probes.py --seeds 20 --out data/probes # 960 probes
python scripts/run_sweep.py --probes data/probes --out data/sweep # GPU, ~4h
python scripts/paper_facts.py --sweep data/sweep/sweep.json # derived claims
python scripts/render_figures.py --sweep data/sweep/sweep.json --figs docs/figures
python scripts/render_tokenbudget_paper.py # this report (HTML + PDF)
python -m tokenbudget ui --port 8000 # interactive console
```

`data/sweep/sweep.json` is the single source of truth: every number in this report traces to that table, and `scripts/paper_facts.py` re-derives the claims from it directly.

References

1. Yao Y., Yu T., Zhang A., Wang C., et al. MiniCPM-V: A GPT-4V Level MLLM on Your Phone. *arXiv:2408.01800*, 2024.
2. OpenBMB. MiniCPM-V 4.6 technical report (2026), native 4x/16x vision-token downsample modes. *openbmb/MiniCPM-V-4.6*.
3. Li Z., Cao S., et al. SigLIP 2: Multilingual Vision-Language Encoders with Improved Semantic Understanding, Localization, and Dense Features. *arXiv:2502.14786*, 2025.
4. Li W., Yuan Y., Liu J., et al. TokenPacker: Efficient Visual Projector for Multimodal LLM. *arXiv:2407.02392*, IJCV 2025.
5. Chen L., Zhao H., Liu T., et al. An Image is Worth 1/2 Tokens After Layer 2: Plug-and-Play Inference Acceleration for Large Vision-Language Models. *ECCV 2024*.
6. Liang Y., Ge C., Tong Z., et al. Not All Patches are What You Need: Expediting Vision Transformers via Token Reorganizations. *ICLR 2022*.
7. Xu R., Ye Y., Yan Y., et al. LLaVA-UHD: An LMM Perceiving Any Aspect Ratio and High-Resolution Images. *ECCV 2024*.
8. Wang P., Bai S., et al. Qwen2-VL: Enhancing Vision-Language Model's Perception of the World at Any Resolution. *arXiv:2409.12191*, 2024.
9. Xu Z., et al. Efficient Multimodal Large Language Models: A Survey. *arXiv:2405.10739*, 2024.
10. Zhai X., Mustafa B., Kolesnikov A., et al. Sigmoid Loss for Language Image Pre-Training (SigLIP). *arXiv:2303.15343*, ICCV 2023.